

L-BFGS-B

3.0

Generated by Doxygen 1.8.13

Contents

1	L-BFGS-B	2
2	File Index	3
2.1	File List	3
3	File Documentation	5
3.1	active.f File Reference	5
3.1.1	Function/Subroutine Documentation	5
3.2	bmvc.f File Reference	6
3.2.1	Function/Subroutine Documentation	6
3.3	cauchy.f File Reference	7
3.3.1	Function/Subroutine Documentation	7
3.4	cmprlb.f File Reference	11
3.4.1	Function/Subroutine Documentation	11
3.5	daxpy.f File Reference	12
3.5.1	Function/Subroutine Documentation	13
3.6	dcopy.f File Reference	13
3.6.1	Function/Subroutine Documentation	13
3.7	dcsrch.f File Reference	14
3.7.1	Function/Subroutine Documentation	15
3.8	dcstep.f File Reference	17
3.8.1	Function/Subroutine Documentation	17
3.9	ddot.f File Reference	18
3.9.1	Function/Subroutine Documentation	18
3.10	dpofa.f File Reference	19
3.10.1	Function/Subroutine Documentation	19
3.11	driver1.f File Reference	20
3.12	driver1.f90 File Reference	20
3.13	driver2.f File Reference	21
3.14	driver2.f90 File Reference	21

3.15 driver3.f File Reference	21
3.16 driver3.f90 File Reference	21
3.17 dscal.f File Reference	21
3.17.1 Function/Subroutine Documentation	21
3.18 dtrsl.f File Reference	22
3.18.1 Function/Subroutine Documentation	22
3.19 errclb.f File Reference	24
3.19.1 Function/Subroutine Documentation	24
3.20 formk.f File Reference	25
3.20.1 Function/Subroutine Documentation	25
3.21 formt.f File Reference	27
3.21.1 Function/Subroutine Documentation	27
3.22 freev.f File Reference	28
3.22.1 Function/Subroutine Documentation	28
3.23 hpsolb.f File Reference	30
3.23.1 Function/Subroutine Documentation	30
3.24 Insrlb.f File Reference	30
3.24.1 Function/Subroutine Documentation	31
3.25 mainlb.f File Reference	33
3.25.1 Function/Subroutine Documentation	33
3.26 matupd.f File Reference	38
3.26.1 Function/Subroutine Documentation	38
3.27 prn1lb.f File Reference	40
3.27.1 Function/Subroutine Documentation	40
3.28 prn2lb.f File Reference	41
3.28.1 Function/Subroutine Documentation	41
3.29 prn3lb.f File Reference	42
3.29.1 Function/Subroutine Documentation	43
3.30 projgr.f File Reference	44
3.30.1 Function/Subroutine Documentation	44
3.31 setulb.f File Reference	45
3.31.1 Function/Subroutine Documentation	45
3.32 subsm.f File Reference	49
3.32.1 Function/Subroutine Documentation	49
3.33 timer.f File Reference	53
3.33.1 Function/Subroutine Documentation	53

1 L-BFGS-B

Software for Large-scale Bound-constrained Optimization

L-BFGS-B is a limited-memory quasi-Newton code for bound-constrained optimization, i.e., for problems where the only constraints are of the form $l \leq x \leq u$. It is intended for problems in which information on the Hessian matrix is difficult to obtain, or for large dense problems. **L-BFGS-B** can also be used for unconstrained problems, and in this case performs similarly to its predecessor, algorithm **L-BFGS** (Harwell routine VA15). The algorithm is implemented in Fortran 77.

Authors

- Ciyou Zhu
- Richard Byrd
- Jorge Nocedal
- Jose Luis Morales

Related Publications

- R. H. Byrd, P. Lu, J. Nocedal and C. Zhu. **A Limited Memory Algorithm for Bound Constrained Optimization** (1995), SIAM Journal on Scientific and Statistical Computing, Vol. 16, Num. 5, pp. 1190-1208
- C. Zhu, R. H. Byrd and J. Nocedal. **L-BFGS-B: Algorithm 778: L-BFGS-B, FORTRAN routines for large scale bound constrained optimization** (1997), ACM Transactions on Mathematical Software, Vol. 23, Num. 4, pp. 550-560
- J.L. Morales and J. Nocedal. **L-BFGS-B: Remark on Algorithm 778: L-BFGS-B, FORTRAN routines for large scale bound constrained optimization** (2011), ACM Transactions on Mathematical Software, Vol. 38, Num. 1
- R. H. Byrd, J. Nocedal and R. B. Schnabel. **Representations of quasi-Newton matrices and their use in limited memory methods** (1994), Mathematical Programming, Vol. 63, pp. 129-156

Note that the subspace minimization in the **LBFGSpp** implementation is an exact minimization subject to the bounds, based on the **BOXCQP** algorithm:

- C. Voglis and I. E. Lagaris, **BOXCQP: An Algorithm for Bound Constrained Convex Quadratic Problems** (2004), 1st International Conference "From Scientific Computing to Computational Engineering", Athens, Greece

For an eagle-eye overview of **L-BFGS-B** and the genealogy **BFGS**->**L-BFGS**->**L-BFGS-B**, see **Henao's Master's thesis**.

Related Software

- [wilmerhenao/L-BFGS-B-NS](#): An L-BFGS-B-NS Optimizer for Non-Smooth Functions
- [pcarbo/lbfgsb-matlab](#): A MATLAB interface for L-BFGS-B
- [bgranzow/L-BFGS-B](#): A pure Matlab implementation of L-BFGS-B (LBFGSB)
- [constantino-garcia/lbfgsb_cpp_wrapper](#): A simple C++ wrapper around the original Fortran L-BFGS-B routine
- [yixuan/LBFGSpp](#): A header-only C++ library for L-BFGS and L-BFGS-B algorithms
- [chokkan/liblbfgs](#): libLBFGS: a library of Limited-memory Broyden-Fletcher-Goldfarb-Shanno (L-BFGS)
- [mkobos/lbfgsb_wrapper](#): Java wrapper for the Fortran L-BFGS-B algorithm
- [yuhonglin/Lbfgsb.jl](#): A Julia wrapper of the l-bfgs-b fortran library
- [Gnimuc/LBFGSB.jl](#): Julia wrapper for L-BFGS-B Nonlinear Optimization Code
- [afbarnard/go-lbfgsb](#): L-BFGS-B optimization for Go, C, Fortran 2003
- [nepluno/lbfgsb-gpu](#): An open source library for the GPU-implementation of L-BFGS-B algorithm
- [Chris00/L-BFGS-ocaml](#): OCaml bindings for L-BFGS
- [dwicke/L-BFGS-B-Lua](#): L-BFGS-B lua wrapper around a L-BFGS-B C implementation
- [avieira/python_lbfgsb](#): Pure Python-based L-BFGS-B implementation
- [ybyygu/rust-lbfgsb](#): Ergonomic bindings to L-BFGS-B code for Rust
- [lbfgsb3c](#): Limited Memory BFGS Minimizer with Bounds on Parameters with optim() 'C' Interface for R

Notes on this repository

I ([J. Schilling](#)) took the freedom to

- put the L-BFGS-B code obtained from [the original website](#) up in this repository,
- divide the subroutines and functions into separate files,
- convert parts of the documentation into a format understandable to [doxygen](#) and
- adjust the `Makefile` to accomodate the separate files and additionally generate a statically linked `liblbfgsb.a` library.

The current release is version 3.0. The distribution file was last changed on 02/08/11.

This work was in no way intending to infringe any copyrights or take credit for others' work. Feel free to contact me at any time in case you noticed something against the rules. Above documentation is obtained from the archived version of the [original manual](#).

2 File Index

2.1 File List

Here is a list of all documented files with brief descriptions:

active.f	5
bmv.f	6
cauchy.f	7
cmprlb.f	11
daxpy.f	12
dcopy.f	13
dcsrch.f	14
dcstep.f	17
ddot.f	18
dpofa.f	19
driver1.f	20
driver1.f90	20
driver2.f	21
driver2.f90	21
driver3.f	21
driver3.f90	21
dscal.f	21
dtrsl.f	22
errclb.f	24
formk.f	25
format.f	27
freev.f	28
hpsolb.f	30
lnsrlb.f	30
mainlb.f	33
matupd.f	38
prn1lb.f	40
prn2lb.f	41
prn3lb.f	42
projgr.f	44
setulb.f	45
subsm.f	49

3 File Documentation

3.1 active.f File Reference

Functions/Subroutines

- subroutine [active](#) (*n*, *l*, *u*, *nbd*, *x*, *iwhere*, *iprint*, *prjctd*, *cnstnd*, *boxed*)

This subroutine initializes iwhere and projects the initial x to the feasible set if necessary.

3.1.1 Function/Subroutine Documentation

3.1.1.1 active()

```
subroutine active (
    integer n,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision, dimension(n) x,
    integer, dimension(n) iwhere,
    integer iprint,
    logical prjctd,
    logical cnstnd,
    logical boxed )
```


This subroutine initializes iwhere and projects the initial x to the feasible set if necessary.

Parameters

<i>n</i>	number of parameters
<i>l</i>	lower bounds on parameters
<i>u</i>	upper bounds on parameters
<i>nbd</i>	indicates which bounds are present
<i>x</i>	position
<i>iwhere</i>	On entry iwhere is unspecified. On exit: iwhere(i)= <ul style="list-style-type: none"> • -1 if x(i) has no bounds • 3 if l(i)=u(i), • 0 otherwise. In cauchy, iwhere is given finer gradations.
<i>iprint</i>	console output flag
<i>prjctd</i>	TODO
<i>cnstnd</i>	TODO
<i>boxed</i>	TODO

Definition at line 34 of file active.f.

Here is the caller graph for this function:



3.2 bmv.f File Reference

Functions/Subroutines

- subroutine **bmv** (*m*, *sy*, *wt*, *col*, *v*, *p*, *info*)

This subroutine computes the product of the 2m x 2m middle matrix in the compact L-BFGS formula of B and a 2m vector v.

3.2.1 Function/Subroutine Documentation

3.2.1.1 bmv()

```

subroutine bmv (
    integer m,
    double precision, dimension(m, m) sy,
    double precision, dimension(m, m) wt,
    integer col,
    double precision, dimension(2*col) v,
    double precision, dimension(2*col) p,
    integer info )
  
```

This subroutine computes the product of the 2m x 2m middle matrix in the compact L-BFGS formula of B and a 2m vector v; it returns the product in p.

Parameters


<i>m</i>	On entry m is the maximum number of variable metric corrections used to define the limited memory matrix. On exit m is unchanged.
<i>sy</i>	On entry sy specifies the matrix S ^t Y. On exit sy is unchanged.
<i>wt</i>	On entry wt specifies the upper triangular matrix J ^t which is the Cholesky factor of (thetaS ^t S+LD ⁻¹ L'). On exit wt is unchanged.
<i>col</i>	On entry col specifies the number of s-vectors (or y-vectors) stored in the compact L-BFGS formula. On exit col is unchanged.
<i>v</i>	On entry v specifies vector v. On exit v is unchanged.

Parameters


<i>p</i>	On entry <i>p</i> is unspecified. On exit <i>p</i> is the product <i>Mv</i> .
<i>info</i>	On entry <i>info</i> is unspecified. On exit <i>info</i> = <ul style="list-style-type: none"> • 0 for normal return, • nonzero for abnormal return when the system to be solved by <i>dtrsl</i> is singular.

Definition at line 36 of file *bmv.f*.

Here is the call graph for this function:



Here is the caller graph for this function:



3.3 cauchy.f File Reference

Functions/Subroutines

- subroutine [cauchy](#) (*n*, *x*, *l*, *u*, *nbd*, *g*, *iorder*, *iwhere*, *t*, *d*, *xcp*, *m*, *wy*, *ws*, *sy*, *wt*, *theta*, *col*, *head*, *p*, *c*, *wbp*, *v*, *nseg*, *iprint*, *sbgnrm*, *info*, *epsmch*)

Compute the Generalized Cauchy Point along the projected gradient direction.

3.3.1 Function/Subroutine Documentation

3.3.1.1 `cauchy()`

```

subroutine cauchy (
    integer n,
    double precision, dimension(n) x,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision, dimension(n) g,
    integer, dimension(n) iorder,
    integer, dimension(n) iwhere,
    double precision, dimension(n) t,
    double precision, dimension(n) d,
    double precision, dimension(n) xcp,
    integer m,
    double precision, dimension(n, col) wy,
    double precision, dimension(n, col) ws,
    double precision, dimension(m, m) sy,
    double precision, dimension(m, m) wt,
    double precision theta,
    integer col,
    integer head,
    double precision, dimension(2*m) p,
    double precision, dimension(2*m) c,
    double precision, dimension(2*m) wbp,
    double precision, dimension(2*m) v,
    integer nseg,
    integer iprint,
    double precision sbgnrm,
    integer info,
    double precision epsmch )

```

For given x , l , u , g (with $sbgnrm > 0$), and a limited memory BFGS matrix B defined in terms of matrices WY , WS , WT , and scalars $head$, col , and $theta$, this subroutine computes the generalized Cauchy point (GCP), defined as the first local minimizer of the quadratic

$$Q(x + s) = g's + 1/2 s'Bs$$

along the projected gradient direction $P(x-tg, l, u)$. The routine returns the GCP in xcp .

Parameters

n	On entry n is the dimension of the problem. On exit n is unchanged.
x	On entry x is the starting point for the GCP computation. On exit x is unchanged.
l	On entry l is the lower bound of x . On exit l is unchanged.
u	On entry u is the upper bound of x . On exit u is unchanged.
nbd	On entry nbd represents the type of bounds imposed on the variables, and must be specified as follows: $nbd(i)=$ <ul style="list-style-type: none"> • 0 if $x(i)$ is unbounded, • 1 if $x(i)$ has only a lower bound, • 2 if $x(i)$ has both lower and upper bounds, and • 3 if $x(i)$ has only an upper bound.
	On exit nbd is unchanged.

Parameters


<i>g</i>	On entry <i>g</i> is the gradient of $f(x)$. <i>g</i> must be a nonzero vector. On exit <i>g</i> is unchanged.
<i>iorder</i>	<i>iorder</i> will be used to store the breakpoints in the piecewise linear path and free variables encountered. On exit, <ul style="list-style-type: none"> • <i>iorder</i>(1),...,<i>iorder</i>(nleft) are indices of breakpoints which have not been encountered; • <i>iorder</i>(nleft+1),...,<i>iorder</i>(nbreak) are indices of encountered breakpoints; and • <i>iorder</i>(nfree),...,<i>iorder</i>(n) are indices of variables which have no bound constraints along the search direction.
<i>iwhere</i>	On entry <i>iwhere</i> indicates only the permanently fixed (<i>iwhere</i> =3) or free (<i>iwhere</i> = -1) components of x . On exit <i>iwhere</i> records the status of the current x variables. <i>iwhere</i> (i)= <ul style="list-style-type: none"> • -3 if $x(i)$ is free and has bounds, but is not moved • 0 if $x(i)$ is free and has bounds, and is moved • 1 if $x(i)$ is fixed at $l(i)$, and $l(i) \neq u(i)$ • 2 if $x(i)$ is fixed at $u(i)$, and $u(i) \neq l(i)$ • 3 if $x(i)$ is always fixed, i.e., $u(i)=x(i)=l(i)$ • -1 if $x(i)$ is always free, i.e., it has no bounds.
<i>t</i>	working array; will be used to store the break points.
<i>d</i>	the Cauchy direction $P(x-tg)-x$
<i>xcp</i>	returns the GCP on exit
<i>m</i>	On entry <i>m</i> is the maximum number of variable metric corrections used to define the limited memory matrix. On exit <i>m</i> is unchanged.
<i>ws</i>	On entry this stores S , a set of s -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wy</i>	On entry this stores Y , a set of y -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>sy</i>	On entry this stores $S'Y$, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wt</i>	On entry this stores the Cholesky factorization of $(\theta S'S + L D^{-1} L')$, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>theta</i>	On entry <i>theta</i> is the scaling factor specifying $B_0 = \theta I$. On exit <i>theta</i> is unchanged.
<i>col</i>	On entry <i>col</i> is the actual number of variable metric corrections stored so far. On exit <i>col</i> is unchanged.
<i>head</i>	On entry <i>head</i> is the location of the first s -vector (or y -vector) in S (or Y). On exit <i>col</i> is unchanged.
<i>p</i>	will be used to store the vector $p = W^\wedge(T)d$.
<i>c</i>	will be used to store the vector $c = W^\wedge(T)(xcp-x)$.
<i>wbp</i>	will be used to store the row of W corresponding to a breakpoint.
<i>v</i>	working array
<i>nseg</i>	On exit <i>nseg</i> records the number of quadratic segments explored in searching for the GCP.

Parameters


<i>iprint</i>	<p>variable that must be set by the user. It controls the frequency and type of output generated:</p> <ul style="list-style-type: none"> • $iprint < 0$ no output is generated; • $iprint = 0$ print only one line at the last iteration; • $0 < iprint < 99$ print also f and $proj\ g$ every $iprint$ iterations; • $iprint = 99$ print details of every iteration except n-vectors; • $iprint = 100$ print also the changes of active set and final x; • $iprint > 100$ print details of every iteration including x and g; <p>When $iprint > 0$, the file <code>iterate.dat</code> will be created to summarize the iteration.</p>
<i>sbgnrm</i>	<p>On entry <code>sbgnrm</code> is the norm of the projected gradient at x. On exit <code>sbgnrm</code> is unchanged.</p>
<i>info</i>	<p>On entry <code>info</code> is 0. On exit <code>info</code> =</p> <ul style="list-style-type: none"> • 0 for normal return, • = nonzero for abnormal return when the the system used in routine <code>bm</code> is singular.
<i>epsrch</i>	machine precision epsilon

Definition at line 133 of file `cauchy.f`.

Here is the call graph for this function:



Here is the caller graph for this function:



3.4 cmprlb.f File Reference

Functions/Subroutines

- subroutine [cmprlb](#) (n, m, x, g, ws, wy, sy, wt, z, r, wa, index, theta, col, head, nfree, cnstnd, info)
This subroutine computes $r = -Z'B(xcp-xk) - Z'g$ by using $wa(2m+1) = W'(xcp-x)$ from subroutine [cauchy](#).

3.4.1 Function/Subroutine Documentation

3.4.1.1 cmprlb()

```

subroutine cmprlb (
    integer n,
    integer m,
    double precision, dimension(n) x,
    double precision, dimension(n) g,
    double precision, dimension(n, m) ws,
    double precision, dimension(n, m) wy,
    double precision, dimension(m, m) sy,
    double precision, dimension(m, m) wt,
    double precision, dimension(n) z,
    double precision, dimension(n) r,
    double precision, dimension(4*m) wa,
    integer, dimension(n) index,
    double precision theta,
    integer col,
    integer head,
    integer nfree,
    logical cnstnd,
    integer info )
  
```

This subroutine computes $r = -Z'B(xcp-xk) - Z'g$ by using $wa(2m+1) = W'(xcp-x)$ from subroutine [cauchy](#).

Parameters

<i>n</i>	number of parameters
<i>m</i>	history size of Hessian approximation
<i>x</i>	position
<i>g</i>	gradient

Parameters


<i>ws</i>	part of L-BFGS matrix
<i>wy</i>	part of L-BFGS matrix
<i>sy</i>	part of L-BFGS matrix
<i>wt</i>	part of L-BFGS matrix
<i>z</i>	TODO
<i>r</i>	TODO
<i>wa</i>	TODO
<i>index</i>	TODO
<i>theta</i>	TODO
<i>col</i>	TODO
<i>head</i>	TODO
<i>nfree</i>	TODO
<i>cnstnd</i>	TODO
<i>info</i>	TODO

Definition at line 29 of file cmprlb.f.

Here is the call graph for this function:



Here is the caller graph for this function:



3.5 daxpy.f File Reference

Functions/Subroutines

- subroutine `daxpy` (`n`, `da`, `dx`, `incx`, `dy`, `incy`)
constant times a vector plus a vector.

3.5.1 Function/Subroutine Documentation

3.5.1.1 daxpy()

```

subroutine daxpy (
    integer n,
    double precision da,
    double precision, dimension(*) dx,
    integer incx,
    double precision, dimension(*) dy,
    integer incy )

```


constant times a vector plus a vector. uses unrolled loops for increments equal to one.

Parameters

<i>n</i>	dimensionality of vectors
<i>da</i>	constant for scaling dx
<i>dx</i>	vector to be scaled and added to dy
<i>incx</i>	spacing between elements in dx
<i>dy</i>	target vector to which da*dx gets added element-wise
<i>incy</i>	spacing between elements in dy

Definition at line 15 of file daxpy.f.

Here is the caller graph for this function:



3.6 dcopy.f File Reference

Functions/Subroutines

- subroutine [dcopy](#) (*n*, *dx*, *incx*, *dy*, *incy*)
copies a vector, *x*, to a vector, *y*.

3.6.1 Function/Subroutine Documentation

3.6.1.1 dcopy()

```

subroutine dcopy (
    integer n,
    double precision, dimension(*) dx,
    integer incx,
    double precision, dimension(*) dy,
    integer incy )

```


copies a vector, x, to a vector, y. uses unrolled loops for increments equal to one.

Parameters

<i>n</i>	dimensionality of vectors
<i>dx</i>	source vector
<i>incx</i>	spacing of elements in dx
<i>dy</i>	target vector
<i>incy</i>	spacing of elements in dy

Definition at line 14 of file dcopy.f.

Here is the caller graph for this function:



3.7 dcsrch.f File Reference

Functions/Subroutines

- subroutine [dcsrch](#) (f, g, stp, ftol, gtol, xtol, stpmin, stpmax, task, isave, dsave)

This subroutine finds a step that satisfies a sufficient decrease condition and a curvature condition.

3.7.1 Function/Subroutine Documentation

3.7.1.1 dcsrch()

```

subroutine dcsrch (
    double precision f,
    double precision g,
    double precision stp,
    double precision ftol,
    double precision gtol,
    double precision xtol,
    double precision stpmin,
    double precision stpmax,
    character(*) task,
    integer, dimension(2) isave,
    double precision, dimension(13) dsave )

```

This subroutine finds a step that satisfies a sufficient decrease condition and a curvature condition.

Each call of the subroutine updates an interval with endpoints stx and sty. The interval is initially chosen so that it contains a minimizer of the modified function

$$\psi(stp) = f(stp) - f(0) - ftol \cdot stp \cdot f'(0).$$

If $\psi(stp) \leq 0$ and $f'(stp) \geq 0$ for some step, then the interval is chosen so that it contains a minimizer of f .

The algorithm is designed to find a step that satisfies the sufficient decrease condition

$$f(stp) \leq f(0) + ftol \cdot stp \cdot f'(0),$$

and the curvature condition

$$|f'(stp)| \leq gtol \cdot |f'(0)|.$$

If $ftol$ is less than $gtol$ and if, for example, the function is bounded below, then there is always a step which satisfies both conditions.

If no step can be found that satisfies both conditions, then the algorithm stops with a warning. In this case stp only satisfies the sufficient decrease condition.

A typical invocation of `dcsrch` has the following outline:

```

task = 'START'
10 continue
call dcsrch( ... )
if (task .eq. 'FG') then
    evaluate the function and the gradient at stp
goto 10
end if

```

NOTE: The user must not alter work arrays between calls.

Parameters

<i>f</i>	On initial entry <i>f</i> is the value of the function at 0. On subsequent entries <i>f</i> is the value of the function at <i>stp</i> . On exit <i>f</i> is the value of the function at <i>stp</i> .
<i>g</i>	On initial entry <i>g</i> is the derivative of the function at 0. On subsequent entries <i>g</i> is the derivative of the function at <i>stp</i> . On exit <i>g</i> is the derivative of the function at <i>stp</i> .
<i>stp</i>	On entry <i>stp</i> is the current estimate of a satisfactory step. On initial entry, a positive initial estimate must be provided. On exit <i>stp</i> is the current estimate of a satisfactory step if task = 'FG'. If task = 'CONV' then <i>stp</i> satisfies the sufficient decrease and curvature condition.
<i>ftol</i>	On entry <i>ftol</i> specifies a nonnegative tolerance for the sufficient decrease condition. On exit <i>ftol</i> is unchanged.
<i>gtol</i>	On entry <i>gtol</i> specifies a nonnegative tolerance for the curvature condition. On exit <i>gtol</i> is unchanged.
<i>xtol</i>	On entry <i>xtol</i> specifies a nonnegative relative tolerance for an acceptable step. The subroutine exits with a warning if the relative difference between <i>sty</i> and <i>stx</i> is less than <i>xtol</i> . On exit <i>xtol</i> is unchanged.
<i>stpmin</i>	On entry <i>stpmin</i> is a nonnegative lower bound for the step. On exit <i>stpmin</i> is unchanged.
<i>stpmax</i>	On entry <i>stpmax</i> is a nonnegative upper bound for the step. On exit <i>stpmax</i> is unchanged.
<i>task</i>	On initial entry <i>task</i> must be set to 'START'. On exit <i>task</i> indicates the required action: <ul style="list-style-type: none"> • If <i>task</i>(1:2) = 'FG' then evaluate the function and derivative at <i>stp</i> and call <i>dcsrc</i> again. • If <i>task</i>(1:4) = 'CONV' then the search is successful. • If <i>task</i>(1:4) = 'WARN' then the subroutine is not able to satisfy the convergence conditions. The exit value of <i>stp</i> contains the best point found during the search. • If <i>task</i>(1:5) = 'ERROR' then there is an error in the input arguments. On exit with convergence, a warning or an error, the variable <i>task</i> contains additional information.
<i>isave</i>	work array
<i>dsave</i>	work array

Definition at line 96 of file *dcsrc*.f.

Here is the call graph for this function:



Here is the caller graph for this function:



3.8 dcstep.f File Reference

Functions/Subroutines

- subroutine [dcstep](#) (stx, fx, dx, sty, fy, dy, stp, fp, dp, brackt, stpmin, stpmax)

This subroutine computes a safeguarded step for a search procedure and updates an interval that contains a step that satisfies a sufficient decrease and a curvature condition.

3.8.1 Function/Subroutine Documentation

3.8.1.1 dcstep()

```

subroutine dcstep (
    double precision stx,
    double precision fx,
    double precision dx,
    double precision sty,
    double precision fy,
    double precision dy,
    double precision stp,
    double precision fp,
    double precision dp,
    logical brackt,
    double precision stpmin,
    double precision stpmax )
  
```

This subroutine computes a safeguarded step for a search procedure and updates an interval that contains a step that satisfies a sufficient decrease and a curvature condition.

The parameter stx contains the step with the least function value. If brackt is set to .true. then a minimizer has been bracketed in an interval with endpoints stx and sty. The parameter stp contains the current step. The subroutine assumes that if brackt is set to .true. then

```
min(stx,sty) < stp < max(stx,sty),
```


and that the derivative at stx is negative in the direction of the step.

Parameters

<i>stx</i>	On entry <i>stx</i> is the best step obtained so far and is an endpoint of the interval that contains the minimizer. On exit <i>stx</i> is the updated best step.
<i>fx</i>	On entry <i>fx</i> is the function at <i>stx</i> . On exit <i>fx</i> is the function at <i>stx</i> .
<i>dx</i>	On entry <i>dx</i> is the derivative of the function at <i>stx</i> . The derivative must be negative in the direction of the step, that is, <i>dx</i> and <i>stp - stx</i> must have opposite signs. On exit <i>dx</i> is the derivative of the function at <i>stx</i> .
<i>sty</i>	On entry <i>sty</i> is the second endpoint of the interval that contains the minimizer. On exit <i>sty</i> is the updated endpoint of the interval that contains the minimizer.
<i>fy</i>	On entry <i>fy</i> is the function at <i>sty</i> . On exit <i>fy</i> is the function at <i>sty</i> .
<i>dy</i>	On entry <i>dy</i> is the derivative of the function at <i>sty</i> . On exit <i>dy</i> is the derivative of the function at the exit <i>sty</i> .
<i>stp</i>	On entry <i>stp</i> is the current step. If <i>brackt</i> is set to <i>.true.</i> then on input <i>stp</i> must be between <i>stx</i> and <i>sty</i> . On exit <i>stp</i> is a new trial step.
<i>fp</i>	On entry <i>fp</i> is the function at <i>stp</i> . On exit <i>fp</i> is unchanged.
<i>dp</i>	On entry <i>dp</i> is the the derivative of the function at <i>stp</i> . On exit <i>dp</i> is unchanged.
<i>brackt</i>	On entry <i>brackt</i> specifies if a minimizer has been bracketed. Initially <i>brackt</i> must be set to <i>.false.</i> On exit <i>brackt</i> specifies if a minimizer has been bracketed. When a minimizer is bracketed <i>brackt</i> is set to <i>.true.</i>
<i>stpmin</i>	On entry <i>stpmin</i> is a lower bound for the step. On exit <i>stpmin</i> is unchanged.
<i>stpmax</i>	On entry <i>stpmax</i> is an upper bound for the step. On exit <i>stpmax</i> is unchanged.

Definition at line 68 of file *dcstep.f*.

Here is the caller graph for this function:



3.9 ddot.f File Reference

Functions/Subroutines

- double precision function **ddot** (*n*, *dx*, *incx*, *dy*, *incy*)
forms the dot product of two vectors.

3.9.1 Function/Subroutine Documentation

3.9.1.1 ddot()

```
double precision function ddot (
    integer n,
    double precision, dimension(*) dx,
    integer incx,
    double precision, dimension(*) dy,
    integer incy )
```

forms the dot product of two vectors. uses unrolled loops for increments equal to one.

Parameters

<i>n</i>	dimensionality of vectors
<i>dx</i>	first vector
<i>incx</i>	spacing of elements in dx
<i>dy</i>	second vector
<i>incy</i>	spacing of elements in dy

Returns

dot product of dx and dy

Definition at line 15 of file ddot.f.

Here is the caller graph for this function:



3.10 dpofa.f File Reference

Functions/Subroutines

- subroutine [dpofa](#) (a, lda, n, info)
factors a double precision symmetric positive definite matrix.

3.10.1 Function/Subroutine Documentation

3.10.1.1 dpofa()

```
subroutine dpofa (
    double precision, dimension(lda,*) a,
    integer lda,
    integer n,
    integer info )
```

factors a double precision symmetric positive definite matrix.


dpofa is usually called by dpoco, but it can be called directly with a saving in time if rcond is not needed. (time for dpoco) = (1 + 18/n)*(time for dpofa) .

Parameters

<i>a</i>	On entry, this is the symmetric matrix to be factored. Only the diagonal and upper triangle are used. On exit this is an upper triangular matrix <i>r</i> so that $a = \text{trans}(r) * r$ where $\text{trans}(r)$ is the transpose. The strict lower triangle is unaltered. If <i>info</i> .ne. 0 , the factorization is not complete.
<i>lda</i>	On entry, this is the leading dimension of the array <i>a</i> . On exit, this value is unaltered.
<i>n</i>	On entry, this is the order of the matrix <i>a</i> . On exit, this value is unaltered.
<i>info</i>	On exit, this signals success: <ul style="list-style-type: none"> • = 0 for normal return. • = <i>k</i> signals an error condition. the leading minor of order <i>k</i> is not positive definite.

Definition at line 29 of file dpofa.f.

Here is the caller graph for this function:



3.11 driver1.f File Reference

Functions/Subroutines

- program **driver**

3.12 driver1.f90 File Reference

Functions/Subroutines

- program **driver**

3.13 driver2.f File Reference

Functions/Subroutines

- program **driver**

3.14 driver2.f90 File Reference

Functions/Subroutines

- program **driver**

3.15 driver3.f File Reference

Functions/Subroutines

- program **driver**

3.16 driver3.f90 File Reference

Functions/Subroutines

- program **driver**

3.17 dscal.f File Reference

Functions/Subroutines

- subroutine **dscal** (n, da, dx, incx)
scales a vector by a constant.

3.17.1 Function/Subroutine Documentation

3.17.1.1 dscal()

```
subroutine dscal (  
    integer n,  
    double precision da,  
    double precision, dimension(*) dx,  
    integer incx )
```


scales a vector by a constant. uses unrolled loops for increment equal to one.

Parameters

n	dimensionality of the vectors
da	scaling factor to be applied on dx
dx	vector to be scaled
$incx$	spacing of elements in dx

Definition at line 13 of file dscal.f.

Here is the caller graph for this function:



3.18 dtrsl.f File Reference

Functions/Subroutines

- subroutine [dtrsl](#) (t , ldt , n , b , job , $info$)
dtrsl solves triangular systems.

3.18.1 Function/Subroutine Documentation

3.18.1.1 dtrsl()

```

subroutine dtrsl (
    double precision, dimension(ldt,*) t,
    integer ldt,
    integer n,
    double precision, dimension(*) b,
    integer job,
    integer info )

```

`dtrsl` solves systems of the form

$$t * x = b$$

or $\text{trans}(t) * x = b$

where t is a triangular matrix of order n . here $\text{trans}(t)$ denotes the transpose of the matrix t .

Parameters


<i>t</i>	On entry, <i>t</i> contains the matrix of the system. the zero elements of the matrix are not referenced, and the corresponding elements of the array can be used to store other information.
<i>ldt</i>	On entry, <i>ldt</i> is the leading dimension of the array <i>t</i> .
<i>n</i>	On entry, <i>n</i> is the order of the system.
<i>b</i>	On entry, <i>b</i> contains the right hand side of the system. On exit, <i>b</i> contains the solution, if <i>info</i> .eq. 0. otherwise <i>b</i> is unaltered.
<i>job</i>	On entry, <i>job</i> specifies what kind of system is to be solved. if <i>job</i> is <ul style="list-style-type: none"> • 00 solve $t*x=b$, <i>t</i> lower triangular, • 01 solve $t*x=b$, <i>t</i> upper triangular, • 10 solve $\text{trans}(t)*x=b$, <i>t</i> lower triangular, • 11 solve $\text{trans}(t)*x=b$, <i>t</i> upper triangular.
<i>info</i>	On exit, <i>info</i> contains zero if the system is nonsingular. otherwise <i>info</i> contains the index of the first zero diagonal element of <i>t</i> .

Definition at line 33 of file dtrsl.f.

Here is the call graph for this function:



Here is the caller graph for this function:



3.19 errclb.f File Reference

Functions/Subroutines

- subroutine `errclb` (`n`, `m`, `factr`, `l`, `u`, `nbd`, `task`, `info`, `k`)
This subroutine checks the validity of the input data.

3.19.1 Function/Subroutine Documentation

3.19.1.1 `errclb()`

```

subroutine errclb (
    integer n,
    integer m,
    double precision factr,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    character*60 task,
    integer info,
    integer k )

```


This subroutine checks the validity of the input data.

Parameters

<i>n</i>	number of parameters
<i>m</i>	history size of approximated Hessian
<i>factr</i>	convergence criterion on function value
<i>l</i>	lower bounds for parameters
<i>u</i>	upper bounds for parameters
<i>nbd</i>	indicates which bounds are present
<i>task</i>	what task is to be performed
<i>info</i>	error/success indicator
<i>k</i>	index of last errournous parameter

Definition at line 17 of file `errclb.f`.

Here is the caller graph for this function:



3.20 formk.f File Reference

Functions/Subroutines

- subroutine [formk](#) (n, nsub, ind, nenter, ileave, indx2, iupdat, updatd, wn, wn1, m, ws, wy, sy, theta, col, head, info)

Forms the LEL^T factorization of the indefinite matrix K .

3.20.1 Function/Subroutine Documentation

3.20.1.1 formk()

```

subroutine formk (
    integer n,
    integer nsub,
    integer, dimension(n) ind,
    integer nenter,
    integer ileave,
    integer, dimension(n) indx2,
    integer iupdat,
    logical updatd,
    double precision, dimension(2*m, 2*m) wn,
    double precision, dimension(2*m, 2*m) wn1,
    integer m,
    double precision, dimension(n, m) ws,
    double precision, dimension(n, m) wy,
    double precision, dimension(m, m) sy,
    double precision theta,
    integer col,
    integer head,
    integer info )

```

This subroutine forms the LEL^T factorization of the indefinite matrix

$K = [-D -Y'ZZ'Y/\theta L_a -R_z'] [L_a -R_z \theta S'AA'S]$ where $E = [-I \ 0] [\ 0 \ I]$

The matrix K can be shown to be equal to the matrix $M^{[-1]}N$ occurring in section 5.1 of [1], as well as to the matrix $Mbar^{[-1]}Nbar$ in section 5.3.

Parameters


<i>n</i>	On entry n is the dimension of the problem. On exit n is unchanged.
<i>nsub</i>	On entry nsub is the number of subspace variables in free set. On exit nsub is not changed.
<i>ind</i>	On entry ind specifies the indices of subspace variables. On exit ind is unchanged.
<i>nenter</i>	On entry nenter is the number of variables entering the free set. On exit nenter is unchanged.
<i>ileave</i>	On entry indx2(ileave),...,indx2(n) are the variables leaving the free set. On exit ileave is unchanged.

Parameters


<i>indx2</i>	On entry <i>indx2</i> (1),..., <i>indx2</i> (<i>nenter</i>) are the variables entering the free set, while <i>indx2</i> (<i>ileave</i>),..., <i>indx2</i> (<i>n</i>) are the variables leaving the free set. On exit <i>indx2</i> is unchanged.
<i>iupdat</i>	On entry <i>iupdat</i> is the total number of BFGS updates made so far. On exit <i>iupdat</i> is unchanged.
<i>updatd</i>	On entry 'updatd' is true if the L-BFGS matrix is updated. On exit 'updatd' is unchanged.
<i>wn</i>	On entry <i>wn</i> is unspecified. On exit the upper triangle of <i>wn</i> stores the LEL^T factorization of the $2*col \times 2*col$ indefinite matrix $[-D -Y'ZZ'Y/\theta L_a -R_z'] [L_a -R_z \theta S'AA'S]$
<i>wn1</i>	On entry <i>wn1</i> stores the lower triangular part of $[Y' ZZ'Y L_a +R_z'] [L_a +R_z S'AA'S]$ in the previous iteration. On exit <i>wn1</i> stores the corresponding updated matrices. The purpose of <i>wn1</i> is just to store these inner products so they can be easily updated and inserted into <i>wn</i> .
<i>m</i>	On entry <i>m</i> is the maximum number of variable metric corrections used to define the limited memory matrix. On exit <i>m</i> is unchanged.
<i>ws</i>	On entry this stores <i>S</i> , a set of <i>s</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wy</i>	On entry this stores <i>Y</i> , a set of <i>y</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>sy</i>	On entry this stores $S'Y$, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>theta</i>	On entry <i>theta</i> is the scaling factor specifying $B_0 = \theta I$. On exit <i>theta</i> is unchanged.
<i>col</i>	On entry <i>col</i> is the actual number of variable metric corrections stored so far. On exit <i>col</i> is unchanged.
<i>head</i>	On entry <i>head</i> is the location of the first <i>s</i> -vector (or <i>y</i> -vector) in <i>S</i> (or <i>Y</i>). On exit <i>col</i> is unchanged.
<i>info</i>	On entry <i>info</i> is unspecified. On exit <i>info</i> <ul style="list-style-type: none"> • = 0 for normal return; • = -1 when the 1st Cholesky factorization failed; • = -2 when the 2st Cholesky factorization failed.

Definition at line 92 of file formk.f.

Here is the call graph for this function:



Here is the caller graph for this function:



3.21 formt.f File Reference

Functions/Subroutines

- subroutine `formt` (`m`, `wt`, `sy`, `ss`, `col`, `theta`, `info`)
Forms the upper half of the pos. def. and symm. T.

3.21.1 Function/Subroutine Documentation

3.21.1.1 formt()

```

subroutine formt (
    integer m,
    double precision, dimension(m, m) wt,
    double precision, dimension(m, m) sy,
    double precision, dimension(m, m) ss,
    integer col,
    double precision theta,
    integer info )
  
```

This subroutine forms the upper half of the pos. def. and symm. $T = \text{theta} * SS + L * D^{(-1)} * L'$, stores T in the upper triangle of the array `wt`, and performs the Cholesky factorization of T to produce $J * J'$, with J' stored in the upper triangle of `wt`.

Parameters


<i>m</i>	history size of approximated Hessian
<i>wt</i>	part of L-BFGS matrix
<i>sy</i>	part of L-BFGS matrix
<i>ss</i>	part of L-BFGS matrix
<i>col</i>	On entry col is the actual number of variable metric corrections stored so far. On exit col is unchanged.
<i>theta</i>	On entry theta is the scaling factor specifying $B_0 = \text{theta } I$. On exit theta is unchanged.
<i>info</i>	error/success indicator

Definition at line 22 of file `formt.f`.

Here is the call graph for this function:



Here is the caller graph for this function:



3.22 freev.f File Reference

Functions/Subroutines

- subroutine `freev` (*n*, *nfree*, *index*, *nenter*, *ileave*, *indx2*, *iwhere*, *wrk*, *updatd*, *cnstnd*, *iprint*, *iter*)
This subroutine counts the entering and leaving variables when iter > 0, and finds the index set of free and active variables at the GCP.

3.22.1 Function/Subroutine Documentation

3.22.1.1 freev()

```

subroutine freev (
    integer n,
    integer nfree,
    integer, dimension(n) index,
    integer nenter,
    integer ileave,
    integer, dimension(n) indx2,
    integer, dimension(n) iwhere,
    logical wrk,
    logical updatd,
    logical cnstnd,
    integer iprint,
    integer iter )

```


This subroutine counts the entering and leaving variables when $iter > 0$, and finds the index set of free and active variables at the GCP.

Parameters

<i>n</i>	number of parameters
<i>nfree</i>	number of free parameters, i.e., those not at their bounds
<i>index</i>	for $i=1,\dots,nfree$, $index(i)$ are the indices of free variables for $i=nfree+1,\dots,n$, $index(i)$ are the indices of bound variables On entry after the first iteration, $index$ gives the free variables at the previous iteration. On exit it gives the free variables based on the determination in cauchy using the array $iwhere$.
<i>nenter</i>	TODO
<i>ileave</i>	TODO
<i>indx2</i>	On entry $indx2$ is unspecified. On exit with $iter > 0$, $indx2$ indicates which variables have changed status since the previous iteration. For $i=1,\dots,nenter$, $indx2(i)$ have changed from bound to free. For $i=ileave+1,\dots,n$, $indx2(i)$ have changed from free to bound.
<i>iwhere</i>	TODO
<i>wrk</i>	TODO
<i>updatd</i>	TODO
<i>cnstnd</i>	indicating whether bounds are present
<i>iprint</i>	control screen output
<i>iter</i>	TODO

Definition at line 34 of file `freev.f`.

Here is the caller graph for this function:



3.23 hpsolb.f File Reference

Functions/Subroutines

- subroutine [hpsolb](#) (n, t, iorder, iheap)

This subroutine sorts out the least element of t, and puts the remaining elements of t in a heap.

3.23.1 Function/Subroutine Documentation

3.23.1.1 hpsolb()

```
subroutine hpsolb (
    integer n,
    double precision, dimension(n) t,
    integer, dimension(n) iorder,
    integer iheap )
```

Parameters

<i>n</i>	On entry n is the dimension of the arrays t and iorder. On exit n is unchanged.
<i>t</i>	On entry t stores the elements to be sorted. On exit t(n) stores the least elements of t, and t(1) to t(n-1) stores the remaining elements in the form of a heap.
<i>iorder</i>	On entry iorder(i) is the index of t(i). On exit iorder(i) is still the index of t(i), but iorder may be permuted in accordance with t.
<i>iheap</i>	On entry iheap should be set as follows: <ul style="list-style-type: none"> • iheap .eq. 0 if t(1) to t(n) is not in the form of a heap, • iheap .ne. 0 if otherwise. On exit iheap is unchanged.

Definition at line 22 of file hpsolb.f.

Here is the caller graph for this function:



3.24 Insrlb.f File Reference

Functions/Subroutines

- subroutine [lnsr1b](#) (*n*, *l*, *u*, *nbd*, *x*, *f*, *fold*, *gd*, *gdold*, *g*, *d*, *r*, *t*, *z*, *stp*, *dnorm*, *dtd*, *xstep*, *stpmx*, *iter*, *ifun*, *iback*, *nfgv*, *info*, *task*, *boxed*, *cnstnd*, *csave*, *isave*, *dsave*)

This subroutine calls subroutine dcsrch from the Minpack2 library to perform the line search. Subroutine dcsrch is safeguarded so that all trial points lie within the feasible region.

3.24.1 Function/Subroutine Documentation

3.24.1.1 Insr1b()

```

subroutine lnsr1b (
    integer n,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision, dimension(n) x,
    double precision f,
    double precision fold,
    double precision gd,
    double precision gdold,
    double precision, dimension(n) g,
    double precision, dimension(n) d,
    double precision, dimension(n) r,
    double precision, dimension(n) t,
    double precision, dimension(n) z,
    double precision stp,
    double precision dnorm,
    double precision dtd,
    double precision xstep,
    double precision stpmx,
    integer iter,
    integer ifun,
    integer iback,
    integer nfgv,
    integer info,
    character*60 task,
    logical boxed,
    logical cnstnd,
    character*60 csave,
    integer, dimension(2) isave,
    double precision, dimension(13) dsave )

```

Parameters


<i>n</i>	number of parameters
<i>l</i>	lower bounds of parameters
<i>u</i>	upper bounds of parameters

Parameters


<i>nbd</i>	<p>On entry <i>nbd</i> represents the type of bounds imposed on the variables, and must be specified as follows: <i>nbd</i>(<i>i</i>)=</p> <ul style="list-style-type: none"> • 0 if <i>x</i>(<i>i</i>) is unbounded, • 1 if <i>x</i>(<i>i</i>) has only a lower bound, • 2 if <i>x</i>(<i>i</i>) has both lower and upper bounds, and • 3 if <i>x</i>(<i>i</i>) has only an upper bound. <p>On exit <i>nbd</i> is unchanged.</p>
<i>x</i>	position
<i>f</i>	function value at <i>x</i>
<i>fold</i>	TODO
<i>gd</i>	TODO
<i>gdold</i>	TODO
<i>g</i>	gradient of <i>f</i> at <i>x</i>
<i>d</i>	TODO
<i>r</i>	TODO
<i>t</i>	TODO
<i>z</i>	TODO
<i>stp</i>	TODO
<i>dnorm</i>	TODO
<i>dtd</i>	TODO
<i>xstep</i>	TODO
<i>stpmx</i>	TODO
<i>iter</i>	TODO
<i>ifun</i>	TODO
<i>iback</i>	TODO
<i>nfgv</i>	TODO
<i>info</i>	TODO
<i>task</i>	TODO
<i>boxed</i>	TODO
<i>cnstnd</i>	TODO
<i>csave</i>	working array
<i>isave</i>	working array
<i>dsave</i>	working array

Definition at line 47 of file *Insrlb.f*.

Here is the call graph for this function:



Here is the caller graph for this function:



3.25 mainlb.f File Reference

Functions/Subroutines

- subroutine `mainlb` (`n`, `m`, `x`, `l`, `u`, `nbd`, `f`, `g`, `factr`, `pgtol`, `ws`, `wy`, `sy`, `ss`, `wt`, `wn`, `snd`, `z`, `r`, `d`, `t`, `xp`, `wa`, `index`, `iwhere`, `indx2`, `task`, `iprint`, `csave`, `lsave`, `isave`, `dsave`)

This subroutine solves bound constrained optimization problems by using the compact formula of the limited memory BFGS updates.

3.25.1 Function/Subroutine Documentation

3.25.1.1 mainlb()

```

subroutine mainlb (
    integer n,
    integer m,
    double precision, dimension(n) x,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision f,
    double precision, dimension(n) g,
    double precision factr,

```

```

double precision pgtol,
double precision, dimension(n, m) ws,
double precision, dimension(n, m) wy,
double precision, dimension(m, m) sy,
double precision, dimension(m, m) ss,
double precision, dimension(m, m) wt,
double precision, dimension(2*m, 2*m) wn,
double precision, dimension(2*m, 2*m) snd,
double precision, dimension(n) z,
double precision, dimension(n) r,
double precision, dimension(n) d,
double precision, dimension(n) t,
double precision, dimension(n) xp,
double precision, dimension(8*m) wa,
integer, dimension(n) index,
integer, dimension(n) iwhere,
integer, dimension(n) indx2,
character*60 task,
integer iprint,
character*60 csave,
logical, dimension(4) lsave,
integer, dimension(23) isave,
double precision, dimension(29) dsave )

```

This subroutine solves bound constrained optimization problems by using the compact formula of the limited memory BFGS updates.

Parameters

<i>n</i>	On entry <i>n</i> is the number of variables. On exit <i>n</i> is unchanged.
<i>m</i>	On entry <i>m</i> is the maximum number of variable metric corrections allowed in the limited memory matrix. On exit <i>m</i> is unchanged.
<i>x</i>	On entry <i>x</i> is an approximation to the solution. On exit <i>x</i> is the current approximation.
<i>l</i>	On entry <i>l</i> is the lower bound of <i>x</i> . On exit <i>l</i> is unchanged.
<i>u</i>	On entry <i>u</i> is the upper bound of <i>x</i> . On exit <i>u</i> is unchanged.
<i>nbd</i>	On entry <i>nbd</i> represents the type of bounds imposed on the variables, and must be specified as follows: <i>nbd</i> (<i>i</i>)= <ul style="list-style-type: none"> • 0 if <i>x</i>(<i>i</i>) is unbounded, • 1 if <i>x</i>(<i>i</i>) has only a lower bound, • 2 if <i>x</i>(<i>i</i>) has both lower and upper bounds, • 3 if <i>x</i>(<i>i</i>) has only an upper bound. On exit <i>nbd</i> is unchanged.
<i>f</i>	On first entry <i>f</i> is unspecified. On final exit <i>f</i> is the value of the function at <i>x</i> .
<i>g</i>	On first entry <i>g</i> is unspecified. On final exit <i>g</i> is the value of the gradient at <i>x</i> .

Parameters

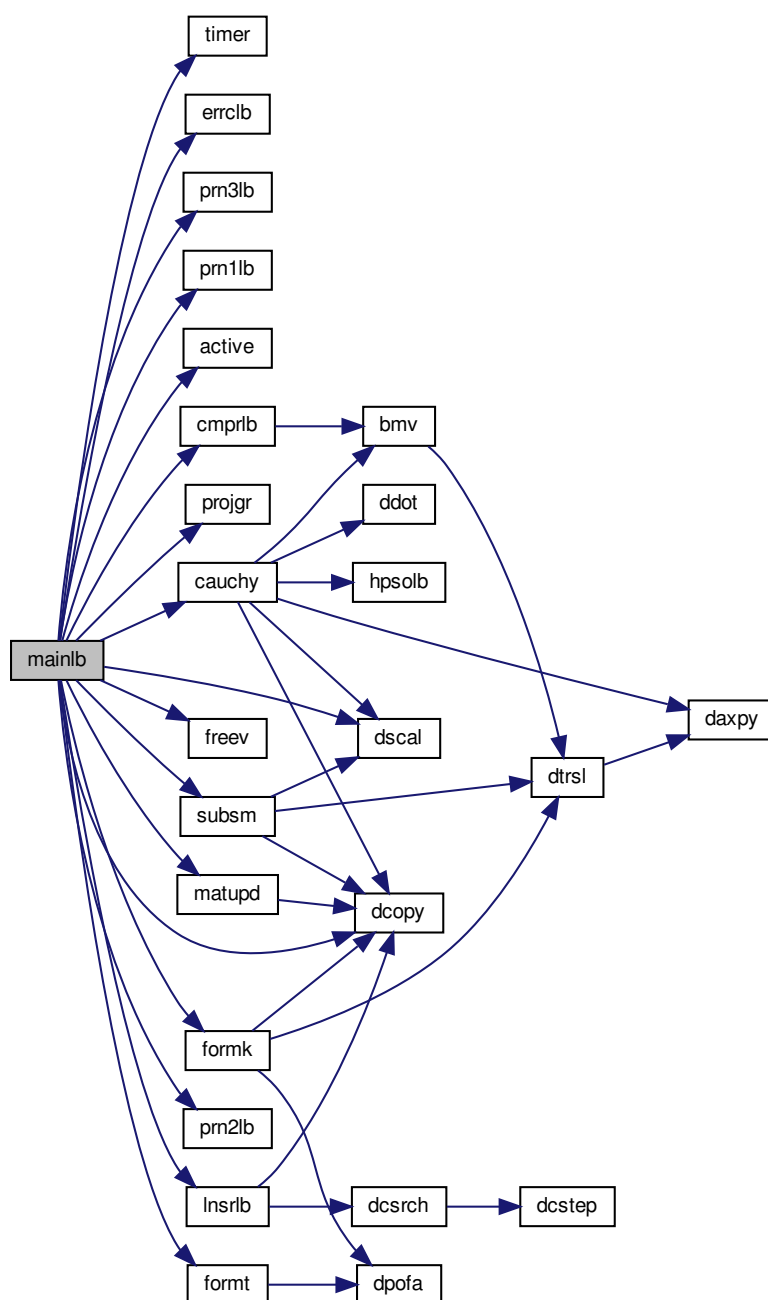
<i>factr</i>	On entry <i>factr</i> ≥ 0 is specified by the user. The iteration will stop when $(f^k - f^{k+1})/\max\{ f^k , f^{k+1} , 1\} \leq \text{factr} \cdot \text{epsmch}$ where <i>epsmch</i> is the machine precision, which is automatically generated by the code. On exit <i>factr</i> is unchanged.
<i>pgtol</i>	On entry <i>pgtol</i> ≥ 0 is specified by the user. The iteration will stop when $\max\{ \text{proj } g_i \mid i = 1, \dots, n\} \leq \text{pgtol}$ where <i>pg_i</i> is the <i>i</i> th component of the projected gradient. On exit <i>pgtol</i> is unchanged.
<i>ws</i>	On entry this stores <i>S</i> , a set of <i>s</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wy</i>	On entry this stores <i>Y</i> , a set of <i>y</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>sy</i>	On entry this stores <i>S</i> ^T <i>Y</i> , that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>ss</i>	On entry this stores <i>S</i> ^T <i>S</i> , that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wt</i>	On entry this stores the Cholesky factorization of $(\theta * S^T S + L D^{(-1)} L^T)$, that defines the limited memory BFGS matrix. See eq. (2.26) in [3]. On exit this array is unchanged.
<i>wn</i>	working array used to store the LEL ^T factorization of the indefinite matrix $K = [-D \ -Y^T Z Z^T Y / \theta \ L_a^T - R_z^T] [L_a \ -R_z \ \theta * S^T A A^T S]$ where $E = [-I \ 0] [\ 0 \ I]$
<i>snd</i>	working array used to store the lower triangular part of $N = [Y^T Z Z^T Y \ L_a^T + R_z^T] [L_a \ +R_z \ S^T A A^T S]$
<i>z</i>	working array used at different times to store the Cauchy point and the Newton point.
<i>r</i>	working array
<i>d</i>	working array
<i>t</i>	working array
<i>xp</i>	working array used to safeguard the projected Newton direction
<i>wa</i>	working array
<i>index</i>	In subroutine <i>freev</i> , <i>index</i> is used to store the free and fixed variables at the Generalized Cauchy Point (GCP).
<i>iwhere</i>	working array used to record the status of the vector <i>x</i> for GCP computation. <i>iwhere</i> (<i>i</i>)= <ul style="list-style-type: none">• 0 or -3 if <i>x</i>(<i>i</i>) is free and has bounds,• 1 if <i>x</i>(<i>i</i>) is fixed at <i>l</i>(<i>i</i>), and <i>l</i>(<i>i</i>) \neq <i>u</i>(<i>i</i>)• 2 if <i>x</i>(<i>i</i>) is fixed at <i>u</i>(<i>i</i>), and <i>u</i>(<i>i</i>) \neq <i>l</i>(<i>i</i>)• 3 if <i>x</i>(<i>i</i>) is always fixed, i.e., <i>u</i>(<i>i</i>)=<i>x</i>(<i>i</i>)=<i>l</i>(<i>i</i>)• -1 if <i>x</i>(<i>i</i>) is always free, i.e., no bounds on it.
<i>indx2</i>	working array Within subroutine <i>cauchy</i> , <i>indx2</i> corresponds to the array <i>iorder</i> . In subroutine <i>freev</i> , a list of variables entering and leaving the free set is stored in <i>indx2</i> , and it is passed on to subroutine <i>formk</i> with this information.
<i>task</i>	working string indicating the current job when entering and leaving this subroutine.

Parameters

<i>iprint</i>	<p>It controls the frequency and type of output generated:</p> <ul style="list-style-type: none">• $iprint < 0$ no output is generated;• $iprint = 0$ print only one line at the last iteration;• $0 < iprint < 99$ print also f and $proj\ g$ every $iprint$ iterations;• $iprint = 99$ print details of every iteration except n-vectors;• $iprint = 100$ print also the changes of active set and final x;• $iprint > 100$ print details of every iteration including x and g; <p>When $iprint > 0$, the file <code>iterate.dat</code> will be created to summarize the iteration.</p>
<i>csave</i>	working string
<i>lsave</i>	working array
<i>isave</i>	working array
<i>dsave</i>	working array

Definition at line 132 of file `mainlb.f`.

Here is the call graph for this function:



Here is the caller graph for this function:



3.26 matupd.f File Reference

Functions/Subroutines

- subroutine [matupd](#) (n, m, ws, wy, sy, ss, d, r, itail, iupdat, col, head, theta, rr, dr, stp, dtd)

This subroutine updates matrices WS and WY, and forms the middle matrix in B.

3.26.1 Function/Subroutine Documentation

3.26.1.1 matupd()

```

subroutine matupd (
    integer n,
    integer m,
    double precision, dimension(n, m) ws,
    double precision, dimension(n, m) wy,
    double precision, dimension(m, m) sy,
    double precision, dimension(m, m) ss,
    double precision, dimension(n) d,
    double precision, dimension(n) r,
    integer itail,
    integer iupdat,
    integer col,
    integer head,
    double precision theta,
    double precision rr,
    double precision dr,
    double precision stp,
    double precision dtd )
  
```

This subroutine updates matrices WS and WY, and forms the middle matrix in B.

Parameters

<i>n</i>	On entry n is the number of variables. On exit n is unchanged.
<i>m</i>	On entry m is the maximum number of variable metric corrections allowed in the limited memory matrix. On exit m is unchanged.

Parameters


<i>ws</i>	On entry this stores S, a set of s-vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wy</i>	On entry this stores Y, a set of y-vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>sy</i>	On entry this stores S'Y, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>ss</i>	On entry this stores S'S, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>d</i>	TODO
<i>r</i>	TODO
<i>itail</i>	TODO
<i>iupdat</i>	TODO
<i>col</i>	On entry col is the actual number of variable metric corrections stored so far. On exit col is unchanged.
<i>head</i>	On entry head is the location of the first s-vector (or y-vector) in S (or Y). On exit col is unchanged.
<i>theta</i>	On entry theta is the scaling factor specifying $B_0 = \theta I$. On exit theta is unchanged.
<i>rr</i>	TODO
<i>dr</i>	TODO
<i>stp</i>	TODO
<i>dtd</i>	TODO

Definition at line 51 of file matupd.f.

Here is the call graph for this function:



Here is the caller graph for this function:



3.27 prn1lb.f File Reference

Functions/Subroutines

- subroutine [prn1lb](#) (n, m, l, u, x, iprint, itfile, epsmch)

This subroutine prints the input data, initial point, upper and lower bounds of each variable, machine precision, as well as the headings of the output.

3.27.1 Function/Subroutine Documentation

3.27.1.1 prn1lb()

```
subroutine prn1lb (
    integer n,
    integer m,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    double precision, dimension(n) x,
    integer iprint,
    integer itfile,
    double precision epsmch )
```


This subroutine prints the input data, initial point, upper and lower bounds of each variable, machine precision, as well as the headings of the output.

Parameters

<i>n</i>	On entry n is the number of variables. On exit n is unchanged.
<i>m</i>	On entry m is the maximum number of variable metric corrections allowed in the limited memory matrix. On exit m is unchanged.
<i>l</i>	On entry l is the lower bound of x. On exit l is unchanged.
<i>u</i>	On entry u is the upper bound of x. On exit u is unchanged.
<i>x</i>	On entry x is an approximation to the solution. On exit x is the current approximation.
<i>iprint</i>	It controls the frequency and type of output generated: <ul style="list-style-type: none"> • <i>iprint</i><0 no output is generated; • <i>iprint</i>=0 print only one line at the last iteration; • 0<<i>iprint</i><99 print also f and proj g every <i>iprint</i> iterations; • <i>iprint</i>=99 print details of every iteration except n-vectors; • <i>iprint</i>=100 print also the changes of active set and final x; • <i>iprint</i>>100 print details of every iteration including x and g; When <i>iprint</i> > 0, the file iterate.dat will be created to summarize the iteration.
<i>itfile</i>	unit number of iterate.dat file
<i>epsmch</i>	machine precision epsilon

Definition at line 40 of file prn1lb.f.

Here is the caller graph for this function:



3.28 prn2lb.f File Reference

Functions/Subroutines

- subroutine [prn2lb](#) (*n*, *x*, *f*, *g*, *iprint*, *itfile*, *iter*, *nfgv*, *nact*, *sbgnrm*, *nseg*, *word*, *iword*, *iback*, *stp*, *xstep*)

This subroutine prints out new information after a successful line search.

3.28.1 Function/Subroutine Documentation

3.28.1.1 prn2lb()

```

subroutine prn2lb (
    integer n,
    double precision, dimension(n) x,
    double precision f,
    double precision, dimension(n) g,
    integer iprint,
    integer itfile,
    integer iter,
    integer nfgv,
    integer nact,
    double precision sbgnrm,
    integer nseg,
    character*3 word,
    integer iword,
    integer iback,
    double precision stp,
    double precision xstep )
  
```

This subroutine prints out new information after a successful line search.

Parameters


<i>n</i>	On entry <i>n</i> is the number of variables. On exit <i>n</i> is unchanged.
<i>x</i>	On entry <i>x</i> is an approximation to the solution. On exit <i>x</i> is the current approximation.

Parameters

<i>f</i>	On first entry <i>f</i> is unspecified. On final exit <i>f</i> is the value of the function at <i>x</i> .
<i>g</i>	On first entry <i>g</i> is unspecified. On final exit <i>g</i> is the value of the gradient at <i>x</i> .
<i>iprint</i>	It controls the frequency and type of output generated: <ul style="list-style-type: none"> • <i>iprint</i><0 no output is generated; • <i>iprint</i>=0 print only one line at the last iteration; • 0<<i>iprint</i><99 print also <i>f</i> and $\text{proj } g$ every <i>iprint</i> iterations; • <i>iprint</i>=99 print details of every iteration except <i>n</i>-vectors; • <i>iprint</i>=100 print also the changes of active set and final <i>x</i>; • <i>iprint</i>>100 print details of every iteration including <i>x</i> and <i>g</i>; When <i>iprint</i> > 0, the file <i>iterate.dat</i> will be created to summarize the iteration.
<i>itfile</i>	unit number of <i>iterate.dat</i> file
<i>iter</i>	TODO
<i>nfgv</i>	TODO
<i>nact</i>	TODO
<i>sbgnrm</i>	TODO
<i>nseg</i>	TODO
<i>word</i>	TODO
<i>iword</i>	TODO
<i>iback</i>	TODO
<i>stp</i>	TODO
<i>xstep</i>	TODO

Definition at line 45 of file *prn2lb.f*.

Here is the caller graph for this function:



3.29 prn3lb.f File Reference

Functions/Subroutines

- subroutine [prn3lb](#) (*n*, *x*, *f*, *task*, *iprint*, *info*, *itfile*, *iter*, *nfgv*, *nintol*, *nskip*, *nact*, *sbgnrm*, *time*, *nseg*, *word*, *iback*, *stp*, *xstep*, *k*, *cachyt*, *sbttime*, *lnscht*)

This subroutine prints out information when either a built-in convergence test is satisfied or when an error message is generated.

3.29.1 Function/Subroutine Documentation

3.29.1.1 prn3lb()

```

subroutine prn3lb (
    integer n,
    double precision, dimension(n) x,
    double precision f,
    character*60 task,
    integer iprint,
    integer info,
    integer itfile,
    integer iter,
    integer nfgv,
    integer nintol,
    integer nskip,
    integer nact,
    double precision sbgnrm,
    double precision time,
    integer nseg,
    character*3 word,
    integer iback,
    double precision stp,
    double precision xstep,
    integer k,
    double precision cachyt,
    double precision sbtime,
    double precision lnscht )

```

This subroutine prints out information when either a built-in convergence test is satisfied or when an error message is generated.

Parameters


<i>n</i>	On entry <i>n</i> is the number of variables. On exit <i>n</i> is unchanged.
<i>x</i>	On entry <i>x</i> is an approximation to the solution. On exit <i>x</i> is the current approximation.
<i>f</i>	On first entry <i>f</i> is unspecified. On final exit <i>f</i> is the value of the function at <i>x</i> .
<i>task</i>	working string indicating the current job when entering and leaving this subroutine.
<i>iprint</i>	TODO
<i>info</i>	TODO
<i>itfile</i>	unit number of iterate.dat file
<i>iter</i>	TODO
<i>nfgv</i>	TODO
<i>nintol</i>	TODO
<i>nskip</i>	TODO
<i>nact</i>	TODO
<i>sbgnrm</i>	TODO
<i>time</i>	TODO
<i>nseg</i>	TODO
<i>word</i>	TODO

Parameters

<i>iback</i>	TODO
<i>stp</i>	TODO
<i>xstep</i>	TODO
<i>k</i>	TODO
<i>cachyt</i>	TODO
<i>sptime</i>	TODO
<i>lnscht</i>	TODO

Definition at line 45 of file prn3lb.f.

Here is the caller graph for this function:



3.30 projgr.f File Reference

Functions/Subroutines

- subroutine [projgr](#) (*n*, *l*, *u*, *nbd*, *x*, *g*, *sbgnrm*)
This subroutine computes the infinity norm of the projected gradient.

3.30.1 Function/Subroutine Documentation

3.30.1.1 projgr()

```

subroutine projgr (
    integer n,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision, dimension(n) x,
    double precision, dimension(n) g,
    double precision sbgnrm )
  
```


This subroutine computes the infinity norm of the projected gradient.

Parameters

<i>n</i>	On entry <i>n</i> is the number of variables. On exit <i>n</i> is unchanged.
<i>l</i>	On entry <i>l</i> is the lower bound of <i>x</i> . On exit <i>l</i> is unchanged.
<i>u</i>	On entry <i>u</i> is the upper bound of <i>x</i> . On exit <i>u</i> is unchanged.
<i>nbd</i>	On entry <i>nbd</i> represents the type of bounds imposed on the variables, and must be specified as follows: <i>nbd</i> (<i>i</i>)= <ul style="list-style-type: none"> • 0 if <i>x</i>(<i>i</i>) is unbounded, • 1 if <i>x</i>(<i>i</i>) has only a lower bound, • 2 if <i>x</i>(<i>i</i>) has both lower and upper bounds, • 3 if <i>x</i>(<i>i</i>) has only an upper bound. On exit <i>nbd</i> is unchanged.
<i>x</i>	On entry <i>x</i> is an approximation to the solution. On exit <i>x</i> is unchanged.
<i>g</i>	On entry <i>g</i> is the gradient. On exit <i>g</i> is unchanged.
<i>sbgnrm</i>	infinity norm of projected gradient

Definition at line 34 of file projgr.f.

Here is the caller graph for this function:



3.31 setulb.f File Reference

Functions/Subroutines

- subroutine [setulb](#) (*n*, *m*, *x*, *l*, *u*, *nbd*, *f*, *g*, *factr*, *pgtol*, *wa*, *iwa*, *task*, *iprint*, *csave*, *lsave*, *isave*, *dsave*)

*This subroutine partitions the working arrays *wa* and *iwa*, and then uses the limited memory BFGS method to solve the bound constrained optimization problem by calling *mainlb*.*

3.31.1 Function/Subroutine Documentation

3.31.1.1 setulb()

```

subroutine setulb (
    integer n,
    integer m,
    double precision, dimension(n) x,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision f,
    double precision, dimension(n) g,
    double precision factr,
    double precision pgtol,
    double precision, dimension(2*m*n + 5*n + 11*m*m + 8*m) wa,
    integer, dimension(3*n) iwa,
    character*60 task,
    integer iprint,
    character*60 csave,
    logical, dimension(4) lsave,
    integer, dimension(44) isave,
    double precision, dimension(29) dsave )

```

This subroutine partitions the working arrays *wa* and *iwa*, and then uses the limited memory BFGS method to solve the bound constrained optimization problem by calling *mainlb*. (The direct method will be used in the subspace minimization.)

Parameters

<i>n</i>	On entry <i>n</i> is the dimension of the problem. On exit <i>n</i> is unchanged.
<i>m</i>	On entry <i>m</i> is the maximum number of variable metric corrections used to define the limited memory matrix. On exit <i>m</i> is unchanged.
<i>x</i>	On entry <i>x</i> is an approximation to the solution. On exit <i>x</i> is the current approximation.
<i>l</i>	On entry <i>l</i> is the lower bound on <i>x</i> . On exit <i>l</i> is unchanged.
<i>u</i>	On entry <i>u</i> is the upper bound on <i>x</i> . On exit <i>u</i> is unchanged.
<i>nbd</i>	On entry <i>nbd</i> represents the type of bounds imposed on the variables, and must be specified as follows: <i>nbd</i> (<i>i</i>)= <ul style="list-style-type: none"> • 0 if <i>x</i>(<i>i</i>) is unbounded, • 1 if <i>x</i>(<i>i</i>) has only a lower bound, • 2 if <i>x</i>(<i>i</i>) has both lower and upper bounds, and • 3 if <i>x</i>(<i>i</i>) has only an upper bound. On exit <i>nbd</i> is unchanged.
<i>f</i>	On first entry <i>f</i> is unspecified. On final exit <i>f</i> is the value of the function at <i>x</i> .
<i>g</i>	On first entry <i>g</i> is unspecified. On final exit <i>g</i> is the value of the gradient at <i>x</i> .

Parameters

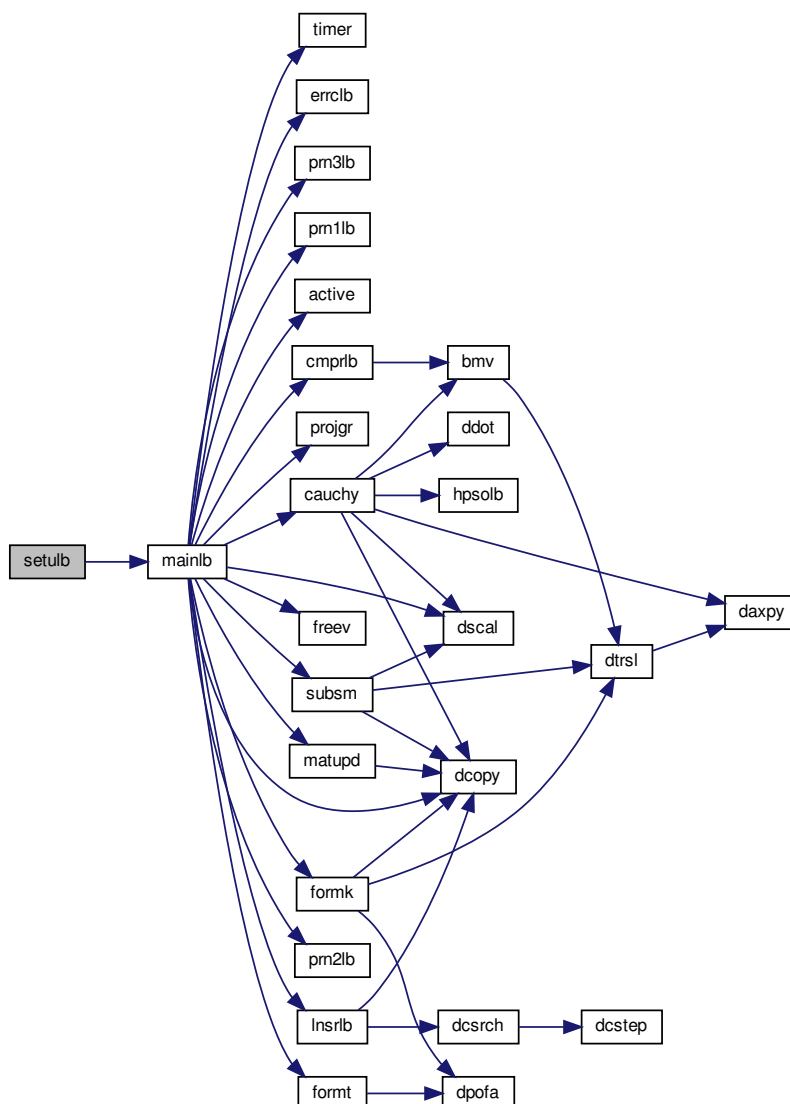
<i>factr</i>	<p>On entry <i>factr</i> ≥ 0 is specified by the user. The iteration will stop when $(f^k - f^{k+1})/\max\{ f^k , f^{k+1} , 1\} \leq \text{factr} \cdot \text{epsmch}$ where <i>epsmch</i> is the machine precision, which is automatically generated by the code. Typical values for <i>factr</i>:</p> <ul style="list-style-type: none"> • 1.d+12 for low accuracy; • 1.d+7 for moderate accuracy; • 1.d+1 for extremely high accuracy. <p>On exit <i>factr</i> is unchanged.</p>
<i>pgtol</i>	<p>On entry <i>pgtol</i> ≥ 0 is specified by the user. The iteration will stop when $\max\{ \text{proj } g_i \mid i = 1, \dots, n\} \leq \text{pgtol}$ where <i>pg_i</i> is the <i>i</i>th component of the projected gradient. On exit <i>pgtol</i> is unchanged.</p>
<i>wa</i>	working array
<i>iwa</i>	working array
<i>task</i>	working string indicating the current job when entering and quitting this subroutine.
<i>iprint</i>	<p>Must be set by the user. It controls the frequency and type of output generated:</p> <ul style="list-style-type: none"> • <i>iprint</i> < 0 no output is generated; • <i>iprint</i> = 0 print only one line at the last iteration; • 0 < <i>iprint</i> < 99 print also <i>f</i> and $\text{proj } g$ every <i>iprint</i> iterations; • <i>iprint</i> = 99 print details of every iteration except <i>n</i>-vectors; • <i>iprint</i> = 100 print also the changes of active set and final <i>x</i>; • <i>iprint</i> > 100 print details of every iteration including <i>x</i> and <i>g</i>; <p>When <i>iprint</i> > 0, the file <i>iterate.dat</i> will be created to summarize the iteration.</p>
<i>csave</i>	working string
<i>lsave</i>	<p>working array; On exit with 'task' = NEW_X, the following information is available:</p> <ul style="list-style-type: none"> • If <i>lsave</i>(1) = .true. then the initial <i>X</i> has been replaced by its projection in the feasible set; • If <i>lsave</i>(2) = .true. then the problem is constrained; • If <i>lsave</i>(3) = .true. then each variable has upper and lower bounds;

Parameters

<i>isave</i>	<p>working array; On exit with 'task' = NEW_X, the following information is available:</p> <ul style="list-style-type: none"> • <i>isave</i>(22) = the total number of intervals explored in the search of Cauchy points; • <i>isave</i>(26) = the total number of skipped BFGS updates before the current iteration; • <i>isave</i>(30) = the number of current iteration; • <i>isave</i>(31) = the total number of BFGS updates prior the current iteration; • <i>isave</i>(33) = the number of intervals explored in the search of Cauchy point in the current iteration; • <i>isave</i>(34) = the total number of function and gradient evaluations; • <i>isave</i>(36) = the number of function value or gradient evaluations in the current iteration; • if <i>isave</i>(37) = 0 then the subspace argmin is within the box; • if <i>isave</i>(37) = 1 then the subspace argmin is beyond the box; • <i>isave</i>(38) = the number of free variables in the current iteration; • <i>isave</i>(39) = the number of active constraints in the current iteration; • $n + 1 - \textit{isave}(40)$ = the number of variables leaving the set of active constraints in the current iteration; • <i>isave</i>(41) = the number of variables entering the set of active constraints in the current iteration.
<i>dsave</i>	<p>working array; On exit with 'task' = NEW_X, the following information is available:</p> <ul style="list-style-type: none"> • <i>dsave</i>(1) = current 'theta' in the BFGS matrix; • <i>dsave</i>(2) = $f(x)$ in the previous iteration; • <i>dsave</i>(3) = $\text{factr} * \text{epsmch}$; • <i>dsave</i>(4) = 2-norm of the line search direction vector; • <i>dsave</i>(5) = the machine precision <i>epsmch</i> generated by the code; • <i>dsave</i>(7) = the accumulated time spent on searching for Cauchy points; • <i>dsave</i>(8) = the accumulated time spent on subspace minimization; • <i>dsave</i>(9) = the accumulated time spent on line search; • <i>dsave</i>(11) = the slope of the line search function at the current point of line search; • <i>dsave</i>(12) = the maximum relative step length imposed in line search; • <i>dsave</i>(13) = the infinity norm of the projected gradient; • <i>dsave</i>(14) = the relative step length in the line search; • <i>dsave</i>(15) = the slope of the line search function at the starting point of the line search; • <i>dsave</i>(16) = the square of the 2-norm of the line search direction vector.

Definition at line 140 of file *setulb.f*.

Here is the call graph for this function:



3.32 subsm.f File Reference

Functions/Subroutines

- subroutine [subsm](#) (n, m, nsub, ind, l, u, nbd, x, d, xp, ws, wy, theta, xx, gg, col, head, iword, wv, wn, iprint, info)

Performs the subspace minimization.

3.32.1 Function/Subroutine Documentation

3.32.1.1 subsm()

```

subroutine subsm (
    integer n,
    integer m,
    integer nsub,
    integer, dimension(nsub) ind,
    double precision, dimension(n) l,
    double precision, dimension(n) u,
    integer, dimension(n) nbd,
    double precision, dimension(n) x,
    double precision, dimension(n) d,
    double precision, dimension(n) xp,
    double precision, dimension(n, m) ws,
    double precision, dimension(n, m) wy,
    double precision theta,
    double precision, dimension(n) xx,
    double precision, dimension(n) gg,
    integer col,
    integer head,
    integer iword,
    double precision, dimension(2*m) wv,
    double precision, dimension(2*m, 2*m) wn,
    integer iprint,
    integer info )

```

Given xcp, l, u, r, an index set that specifies the active set at xcp, and an L-BFGS matrix B (in terms of WY, WS, SY, WT, head, col, and theta), this subroutine computes an approximate solution of the subspace problem

$$(P) \min Q(x) = r'(x-xcp) + 1/2 (x-xcp)' B (x-xcp)$$

subject to $l \leq x \leq u$ $x_i = xcp_i$ for all i in $A(xcp)$

along the subspace unconstrained Newton direction

$$d = -(Z'BZ)^{(-1)} r.$$

The formula for the Newton direction, given the L-BFGS matrix and the Sherman-Morrison formula, is

$$d = (1/\theta)r + (1/\theta^2) Z'WK^{(-1)}W'Z r.$$

$$\text{where } K = [-D -Y'ZZ'Y/\theta L_a -R_z'] [L_a -R_z \theta S'AA'S]$$

Note that this procedure for computing d differs from that described in [1]. One can show that the matrix K is equal to the matrix $M^{[-1]}N$ in that paper.

Parameters

<i>n</i>	On entry n is the dimension of the problem. On exit n is unchanged.
<i>m</i>	On entry m is the maximum number of variable metric corrections used to define the limited memory matrix. On exit m is unchanged.
<i>nsub</i>	On entry nsub is the number of free variables. On exit nsub is unchanged.
<i>ind</i>	On entry ind specifies the coordinate indices of free variables. On exit ind is unchanged.

Parameters


<i>l</i>	On entry <i>l</i> is the lower bound of <i>x</i> . On exit <i>l</i> is unchanged.
<i>u</i>	On entry <i>u</i> is the upper bound of <i>x</i> . On exit <i>u</i> is unchanged.
<i>nbd</i>	On entry <i>nbd</i> represents the type of bounds imposed on the variables, and must be specified as follows: <i>nbd</i> (<i>i</i>)= <ul style="list-style-type: none"> • 0 if <i>x</i>(<i>i</i>) is unbounded, • 1 if <i>x</i>(<i>i</i>) has only a lower bound, • 2 if <i>x</i>(<i>i</i>) has both lower and upper bounds, and • 3 if <i>x</i>(<i>i</i>) has only an upper bound. On exit <i>nbd</i> is unchanged.
<i>x</i>	On entry <i>x</i> specifies the Cauchy point <i>xcp</i> . On exit <i>x</i> (<i>i</i>) is the minimizer of <i>Q</i> over the subspace of free variables.
<i>d</i>	On entry <i>d</i> is the reduced gradient of <i>Q</i> at <i>xcp</i> . On exit <i>d</i> is the Newton direction of <i>Q</i> .
<i>xp</i>	used to safeguard the projected Newton direction
<i>xx</i>	On entry it holds the current iterate. On output it is unchanged.
<i>gg</i>	On entry it holds the gradient at the current iterate. On output it is unchanged.
<i>ws</i>	On entry this stores <i>S</i> , a set of <i>s</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>wy</i>	On entry this stores <i>Y</i> , a set of <i>y</i> -vectors, that defines the limited memory BFGS matrix. On exit this array is unchanged.
<i>theta</i>	On entry <i>theta</i> is the scaling factor specifying $B_0 = \text{theta } I$. On exit <i>theta</i> is unchanged.
<i>col</i>	On entry <i>col</i> is the actual number of variable metric corrections stored so far. On exit <i>col</i> is unchanged.
<i>head</i>	On entry <i>head</i> is the location of the first <i>s</i> -vector (or <i>y</i> -vector) in <i>S</i> (or <i>Y</i>). On exit <i>col</i> is unchanged.
<i>word</i>	On entry <i>word</i> is unspecified. On exit <i>word</i> specifies the status of the subspace solution. <i>word</i> = <ul style="list-style-type: none"> • 0 if the solution is in the box, • 1 if some bound is encountered.
<i>wv</i>	working array
<i>wn</i>	On entry the upper triangle of <i>wn</i> stores the LEL^T factorization of the indefinite matrix $K = [-D -Y'ZZ'Y/\text{theta } L_a -R_z'] [L_a -R_z \text{theta} *S'AA'S]$ where $E = [-I \ 0] [\ 0 \ I]$ On exit <i>wn</i> is unchanged.

Parameters


<i>iprint</i>	<p>must be set by the user; It controls the frequency and type of output generated:</p> <ul style="list-style-type: none"> • $iprint < 0$ no output is generated; • $iprint = 0$ print only one line at the last iteration; • $0 < iprint < 99$ print also f and $proj\ g$ every $iprint$ iterations; • $iprint = 99$ print details of every iteration except n-vectors; • $iprint = 100$ print also the changes of active set and final x; • $iprint > 100$ print details of every iteration including x and g; <p>When $iprint > 0$, the file <code>iterate.dat</code> will be created to summarize the iteration.</p>
<i>info</i>	<p>On entry <code>info</code> is unspecified. On exit <code>info</code> =</p> <ul style="list-style-type: none"> • 0 for normal return, • nonzero for abnormal return when the matrix K is ill-conditioned.

Definition at line 127 of file `subsm.f`.

Here is the call graph for this function:



Here is the caller graph for this function:



3.33 timer.f File Reference

Functions/Subroutines

- subroutine `timer` (`ttime`)

This routine computes cpu time in double precision.

3.33.1 Function/Subroutine Documentation

3.33.1.1 timer()

```
subroutine timer (  
    double precision ttime )
```


This routine computes cpu time in double precision; it makes use of the intrinsic `f90 cpu_time` therefore a conversion type is needed.

Parameters

<code>ttime</code>	CPU time in double precision
--------------------	------------------------------

Definition at line 11 of file `timer.f`.

Here is the caller graph for this function:



Index

active
 active.f, 5
active.f, 5
 active, 5

bmv
 bmv.f, 6
bmv.f, 6
 bmv, 6

cauchy
 cauchy.f, 7
cauchy.f, 7
 cauchy, 7
cmprlb
 cmprlb.f, 11
cmprlb.f, 11
 cmprlb, 11

daxpy
 daxpy.f, 13
daxpy.f, 12
 daxpy, 13
dcopy
 dcopy.f, 13
dcopy.f, 13
 dcopy, 13
dcsrch
 dcsrch.f, 15
dcsrch.f, 14
 dcsrch, 15
dcstep
 dcstep.f, 17
dcstep.f, 17
 dcstep, 17
ddot
 ddot.f, 18
ddot.f, 18
 ddot, 18
dpofa
 dpofa.f, 19
dpofa.f, 19
 dpofa, 19
driver1.f, 20
driver1.f90, 20
driver2.f, 21
driver2.f90, 21
driver3.f, 21
driver3.f90, 21
dscal
 dscal.f, 21
dscal.f, 21
 dscal, 21
dtrsl
 dtrsl.f, 22
dtrsl.f, 22

dtrsl, 22

errclb
 errclb.f, 24
errclb.f, 24
 errclb, 24

formk
 formk.f, 25
formk.f, 25
 formk, 25
formt
 formt.f, 27
formt.f, 27
 formt, 27
freev
 freev.f, 28
freev.f, 28
 freev, 28

hpsolb
 hpsolb.f, 30
hpsolb.f, 30
 hpsolb, 30

lnsrbl
 lnsrbl.f, 31
lnsrbl.f, 30
 lnsrbl, 31

mainlb
 mainlb.f, 33
mainlb.f, 33
 mainlb, 33
matupd
 matupd.f, 38
matupd.f, 38
 matupd, 38

prn1lb
 prn1lb.f, 40
prn1lb.f, 40
 prn1lb, 40
prn2lb
 prn2lb.f, 41
prn2lb.f, 41
 prn2lb, 41
prn3lb
 prn3lb.f, 43
prn3lb.f, 42
 prn3lb, 43
projgr
 projgr.f, 44
projgr.f, 44
 projgr, 44

setulb

- setulb.f, [45](#)
- setulb.f, [45](#)
 - setulb, [45](#)
- subsm
 - subsm.f, [49](#)
- subsm.f, [49](#)
 - subsm, [49](#)
- timer
 - timer.f, [53](#)
- timer.f, [53](#)
 - timer, [53](#)